

Zoo:
A framework for the verification
of concurrent OCaml 5 programs
using separation logic



Clément
Allain



Gabriel
Scherer

Context

Zoo in practice

Zoo features

Physical equality in HeapLang

Physical equality in OCaml

Future work

Context

Verification of *fine-grained concurrent OCaml 5* programs



Saturn

Kcas

Parabs



Iris artillery

- ▶ higher-order ghost state
- ▶ user-defined ghost state
- ▶ invariants
- ▶ atomic updates
- ▶ prophecy variables

HeapLang, the canonical Iris language

- ▶ **simple & expressive, but**
- ▶ **lacks basic abstractions**
 - ▶ tuples, records
 - ▶ algebraic data types (ADTs)
 - ▶ mutually recursive functions
- ▶ **lacks a standard library**
- ▶ **physical equality is problematic**
 - ▶ restricted to “unboxed” values
 - ▶ incompatible with OCaml

Context

Zoo in practice

Zoo features

Physical equality in HeapLang

Physical equality in OCaml

Future work

The big picture



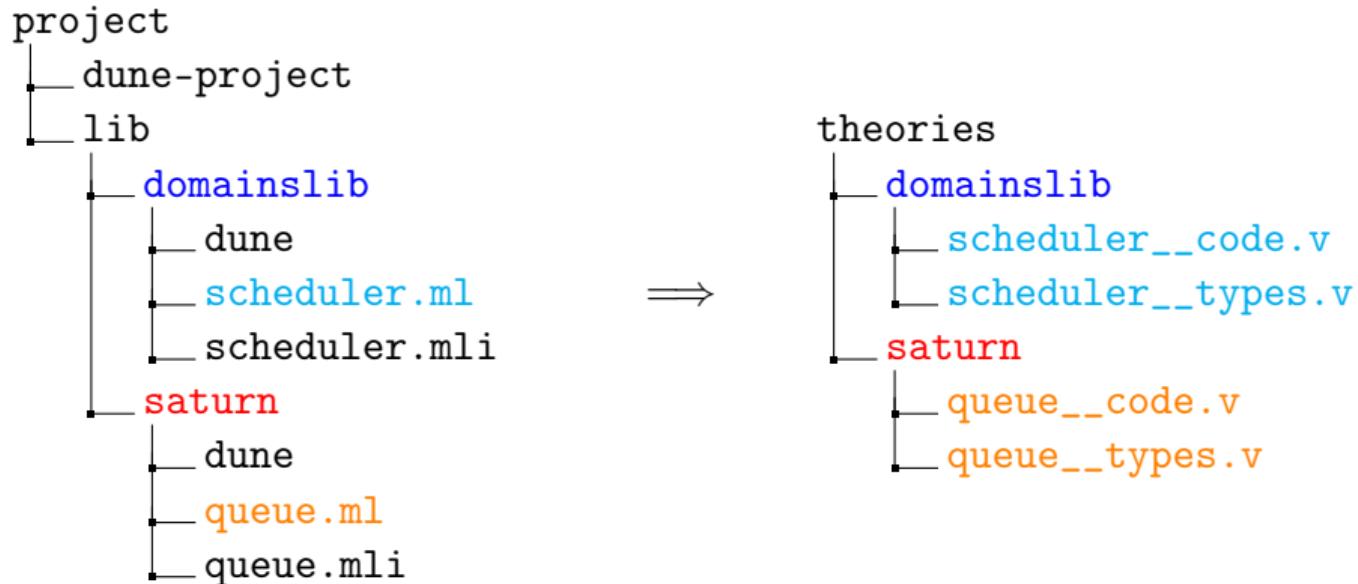
ocaml2zoo
→



Zoo

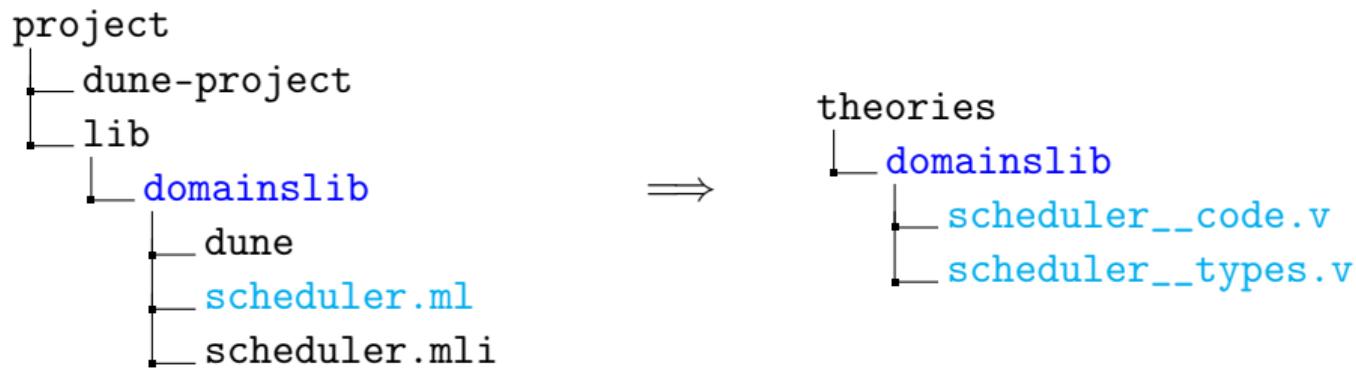


Zoo in practice



```
$ ocaml2zoo project theories
```

Zoo in practice



```
$ ocaml2zoo project theories
```

Zoo in practice

```
Lemma stack_push_spec_seq t  $\iota$  v :  
{{{  
  stack_model t vs  
}}} }  
  stack_push t v  
{{{  
  RET ();  
  stack_model t (v :: vs)  
}}}.  
Proof.  
...  
Qed.
```

```
Lemma stack_push_spec_atomic t  $\iota$  v :  
<<<  
  stack_inv t  $\iota$   
|  $\forall$  vs,  
  stack_model t vs  
>>>  
  stack_push t v @  $\uparrow$  $\iota$   
<<<  
  stack_model t (v :: vs)  
| RET (); True  
>>>.  
Proof.  
...  
Qed.
```

Context

Zoo in practice

Zoo features

Physical equality in HeapLang

Physical equality in OCaml

Future work

Algebraic data types

```
type 'a t =
| Nil
| Cons of 'a * 'a t

let rec map fn t =
  match t with
  | Nil -> Nil
  | Cons (x, t) ->
    let y = fn x in
    Cons (y, map fn t)
```

```
Notation "'Nil'" := (
  in_type "t" 0
)(in custom zoo_tag).

Notation "'Cons'" := (
  in_type "t" 1
)(in custom zoo_tag).
```

```
Definition map : val :=
rec: "map" "fn" "t" =>
  match: "t" with
  | Nil => §Nil
  | Cons "x" "t" =>
    let: "y" := "fn" "x" in
    'Cons( "y", "map" "fn" "t" )
end.
```

Records

```
type 'a t =
{ mutable f1: 'a;
  mutable f2: 'a;
}
```

```
let swap t =
  let f1 = t.f1 in
  t.f1 <- t.f2 ;
  t.f2 <- f1
```

```
Notation "'f1'" := (
  in_type "t" 0
)(in custom zoo_field).
Notation "'f2'" := (
  in_type "t" 1
)(in custom zoo_field).
```

```
Definition swap : val :=
  fun: "t" =>
    let: "f1" := "t".{f1} in
    "t" <-{f1} "t".{f2} ;;
    "t" <-{f2} "f1".
```

Inline records

```
type 'a node =
| Null
| Node of
  { mutable next: 'a node;
    mutable data: 'a;
  }
```

```
Notation "'Null'" := (
  in_type "node" 0
)(in custom zoo_tag).
Notation "'Node'" := (
  in_type "node" 1
)(in custom zoo_tag).

Notation "'next'" := (
  in_type "node.Node" 0
)(in custom zoo_field).
Notation "'data'" := (
  in_type "node.Node" 1
)(in custom zoo_field).
```

Mutually recursive functions

```
Definition f_g := (
  recs: "f" "x" => "g" "x"
  and:  "g" "x" => "f" "x"
)%zoo_recs.

(* boilerplate *)

let rec f x = g x
and g x = f x

Definition f := ValRecs 0 f_g.
Definition g := ValRecs 1 f_g.

Instance : AsValRecs' f 0 f_g [f;g].
Proof. done. Qed.
Instance : AsValRecs' g 1 f_g [f;g].
Proof. done. Qed.
```

Concurrency

`Atomic.set e1 e2`

$e_1 \leftarrow e_2$

`Atomic.exchange e1 e2`

Xchg $e_1.[\text{contents}] e_2$

`Atomic.compare_and_set e1 e2 e3`

CAS $e_1.[\text{contents}] e_2 e_3$

`Atomic.fetch_and_add e1 e2`

FAA $e_1.[\text{contents}] e_2$

`type t = { ...; mutable f: τ [@atomic]; ... }`

`Atomic.Loc.exchange [%atomic.loc e1.f] e2`

Xchg $e_1.[f] e_2$

`Atomic.Loc.compare_and_set [%atomic.loc e1.f] e2 e3`

CAS $e_1.[f] e_2 e_3$

`Atomic.Loc.fetch_and_add [%atomic.loc e1.f] e2`

FAA $e_1.[f] e_2$

Standard library

- ▶ Array
- ▶ Dynarray
- ▶ List
- ▶ Stack
- ▶ Queue
- ▶ Deque
- ▶ Domain
- ▶ Atomic_array
- ▶ Mutex
- ▶ Semaphore
- ▶ Condition
- ▶ Ivar

Diaframe (Ike Mulder *et al.*) support

Proof.

...

```
iInv "Hinv" as "(:inv_inner =1)".
```

...

```
iSplitR ... { iFrame. iSteps. }
```

```
iInv "Hinv" as "(:inv_inner =2)".
```

...

```
iSplitR ... { iFrame. iSteps. }
```

...

Qed.

Context

Zoo in practice

Zoo features

Physical equality in HeapLang

Physical equality in OCaml

Future work

Restriction on physical comparison

```
Definition val_is_unboxed v :=
  match v with
  | LitV lit =>
    lit_is_unboxed lit
  | InjLV (LitV lit) =>
    lit_is_unboxed lit
  | InjRV (LitV lit) =>
    lit_is_unboxed lit
  | _ =>
    False
end.
```

```
Definition vals_compare_safe v1 v2 :=
  val_is_unboxed v1 ∨ val_is_unboxed v2.
```

Treiber stack

```
type 'a t =
  'a list Atomic.t

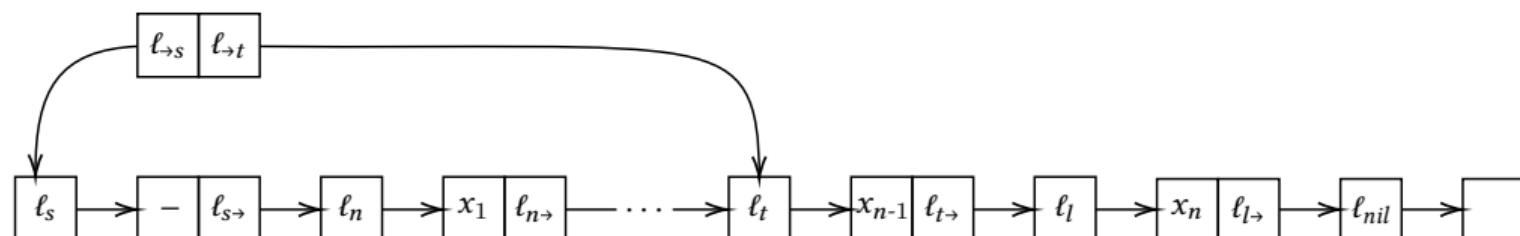
let rec push t v =
  let old = Atomic.get t in
  let new_ = v :: old in
  if not @@ Atomic.compare_and_set t old new_ then
    push t v
```

```
Definition push : val :=
  rec: "push" "t" "v" :=
    let: "old" := ! "t" in
    let: "new" := SOME (ref ("v", "old")) in
    if: CAS "t" "old" "new" then #() else
      "push" "t" "v".
```

Michael-Scott queue

```
type ('a, _) node =
| Null : ('a, [> `Null]) node
| Node : { mutable next: ('a, [`Null | `Node]) node [@atomic];
            mutable data: 'a; }
          -> ('a, [> `Node]) node

type 'a t =
{ mutable front: ('a, [`Node]) node [@atomic];
  mutable back: ('a, [`Node]) node [@atomic]; }
```



RDCSS

`NewRDCSS(n)` \triangleq `ref(inl(n));`

```

rec Get( $\ell_n$ )  $\triangleq$ 
  match ! $\ell_n$  with
    inl( $n$ )  $\Rightarrow n$ 
  | inr( $\ell_{desc}$ )  $\Rightarrow$  Complete( $\ell_{desc}$ ,  $\ell_n$ ); Get( $\ell_n$ )
  end;

```

```

15 Complete( $\ell_{desc}$ ,  $\ell_n$ )  $\triangleq$ 
16 let ( $\ell_m, m_1, n_1, n_2, p$ ) = ! $\ell_{desc}$ ;
17 let tid = NewGhostId;
18 let  $m = !\ell_m$ ;
19 let  $n_{new} = \text{if } m = m_1 \text{ then } n_2 \text{ else } n_1$ ;
20 Resolve(CmpX( $\ell_n$ , inr( $\ell_{desc}$ ), inl( $n_{new}$ )),  $p, tid$ );
21 ()

```

```

1 RDCSS( $\ell_m, \ell_n, m_1, n_1, n_2$ )  $\triangleq$ 
2 let  $p = \text{NewProp}$ ;
3 let  $\ell_{desc} = \text{ref}(\ell_m, m_1, n_1, n_2, p)$ ;
4 rec rdcssinner()  $=$ 
5   let ( $v, b$ ) = CmpX( $\ell_n$ , inl( $n_1$ ), inr( $\ell_{desc}$ ));
6   match  $v$  with
7     inl( $n$ )  $\Rightarrow$ 
8       if  $b$  then
9         Complete( $\ell_{desc}$ ,  $\ell_n$ );  $n_1$ 
10        else  $n$ 
11      | inr( $\ell'_{desc}$ )  $\Rightarrow$ 
12        Complete( $\ell'_{desc}$ ,  $\ell_n$ ); rdcssinner()
13    end;
14  rdcssinner()

```

Physical comparison decides Rocq equality

```
Definition bin_op_eval op v1 v2 :=
  if decide (op = EqOp) then
    if decide (vals_compare_safe v1 v2) then
      Some $ LitV $ LitBool $ bool Decide (v1 = v2)
    else
      None
  else
    . . .
```

Context

Zoo in practice

Zoo features

Physical equality in HeapLang

Physical equality in OCaml

Future work

Classification of Zoo values

- ▶ boolean
- ▶ integer
- ▶ mutable block (pointer)
- ▶ immutable block (tag and fields)
- ▶ function

Non-deterministic semantics

```
let x1 = Some ()  
let x2 = Some ()  
let test1 = x1 == x1 (* true *)  
let test2 = x1 == x2 (* false *)
```

What guarantees when physical equality
(1) returns **true**,
(2) returns **false**?

Sharing

```
let test1 = Some 0 == Some 0 (* true *)
let test2 = [0;1] == [0;1] (* true *)
```

Value representation conflicts

```
let test1 = Obj.repr false == Obj.repr 0 (* true *)
let test2 = Obj.repr None == Obj.repr 0 (* true *)
let test3 = Obj.repr [] == Obj.repr 0 (* true *)
```

Sharing + conflicts

```
type any =
  Any : 'a -> any

let test1 = Any false == Any 0 (* true *)
let test2 = Any None == Any 0 (* true *)
let test3 = Any [] == Any 0 (* true *)
```

Treiber stack

```
type 'a t =
  'a list Atomic.t

let create () =
  Atomic.make []

let rec push t v =
  let old = Atomic.get t in
  let new_ = v :: old in
  if not @@ Atomic.compare_and_set t old new_ then (
    Domain.cpu_relax ();
    push t v
  )
```

Treiber stack specification

```
Lemma stack_push_spec t  $\iota$  v :
```

```
<<<
```

```
  stack_inv t  $\iota$ 
```

```
|  $\forall \forall$  vs,
```

```
  stack_model t vs
```

```
>>>
```

```
  stack_push t v @  $\uparrow \iota$ 
```

```
<<<
```

```
  stack_model t (v :: vs)
```

```
| RET () ; True
```

```
>>>.
```

```
Proof.
```

```
...
```

```
Qed.
```

Unsharing

```
let x = Some 0  
let test = x == x (* false *)
```



Clément Allain
Impossible! Unique identity.



Armaël Guéneau
This would be *unsharing*.



Vincent Lviron
It's possible!

Eio.Rcfd

```
type state = Open of Unix.file_descr | Closing of (unit -> unit)
type t = { mutable ops: int [@atomic]; mutable state: state [@atomic] }

let make fd = { ops= 0; state= Open fd }

let closed = Closing (fun () -> ())
let close t =
  match t.state with
  | Closing _ -> false
  | Open fd as prev ->
    let close () = Unix.close fd in
    let next = Closing close in
    if Atomic.Loc.compare_and_set [%atomic.loc t.state] prev next then
      ...
    else
      false
```

Generative constructors

```
type 'a list =
| Nil
| Cons of 'a * 'a list [@generative]

type state =
| Open of Unix.file_descr [@generative] [@zoo.reveal]
| Closing of (unit -> unit)
```

Context

Zoo in practice

Zoo features

Physical equality in HeapLang

Physical equality in OCaml

Future work

Future work

- ▶ exceptions
- ▶ algebraic effects
- ▶ modules & functors
- ▶ weak memory
- ▶ coupling with semi-automated verification

Thank you for your attention!